

Well-Structured Futures and Cache Locality

MAURICE HERLIHY and ZHIYU LIU, Brown University

In fork-join parallelism, a sequential program is split into a directed acyclic graph of tasks linked by directed dependency edges, and the tasks are executed, possibly in parallel, in an order consistent with their dependencies. A popular and effective way to extend fork-join parallelism is to allow threads to create futures. A thread creates a future to hold the results of a computation, which may or may not be executed in parallel. That result is returned when some thread touches that future, blocking if necessary until the result is ready.

Recent research has shown that although futures can, of course, enhance parallelism in a structured way, they can have a deleterious effect on cache locality. In the worst case, futures can incur $\Omega(PT_\infty + tT_\infty)$ deviations, which implies $\Omega(CPT_\infty + CtT_\infty)$ additional cache misses, where C is the number of cache lines, P is the number of processors, t is the number of touches, and T_∞ is the computation span. Since cache locality has a large impact on software performance on modern multicores, this result is troubling.

In this article, we show that if futures are used in a simple, disciplined way, then the situation is much better: if each future is touched only once, either by the thread that created it or by a later descendant of the thread that created it, then parallel executions with work stealing can incur at most $O(CPT_\infty^2)$ additional cache misses—a substantial improvement. This structured use of futures is characteristic of many (but not all) parallel applications.

Categories and Subject Descriptors: C.4 [Performance of Systems]: Performance Attributes; D.3.3 [Language Constructs and Features]: Concurrent Programming Structures

General Terms: Performance

Additional Key Words and Phrases: Scheduling, work stealing, futures, parallel programming, cache locality, performance bounds

ACM Reference Format:

Maurice Herlihy and Zhiyu Liu. 2016. Well-structured futures and cache locality. ACM Trans. Parallel Comput. 2, 4, Article 22 (February 2016), 20 pages.
DOI: <http://dx.doi.org/10.1145/2858650>

1. INTRODUCTION

Futures [Halstead 1984, 1985] are an attractive way to structure many parallel programs because they are easy to reason about (especially if the futures have no side effects), and they lend themselves well to sophisticated dynamic scheduling algorithms, such as work stealing [Blumofe and Leiserson 1999] and its variations, that ensure high processor utilization. At the same time, however, modern multicore architectures employ complex multilevel memory hierarchies, and technology trends are increasing the relative performance differences among the various levels of memory. As a result, processor utilization can no longer be the sole figure of merit for schedulers. Instead, the *cache locality* of the parallel execution will become increasingly critical to overall

This work is supported by the National Science Foundation under grants 1420673 and 1301924.

Authors' address: M. Herlihy and Z. Liu, Computer Science Department, Brown University, Providence, Rhode Island 02912 USA; emails: mph@cs.brown.edu, zhiyu_liu@brown.edu.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.

© 2016 ACM 2329-4949/2016/02-ART22 \$15.00

DOI: <http://dx.doi.org/10.1145/2858650>

performance. As a result, cache locality will increasingly join processor utilization as a criterion for evaluating dynamic scheduling algorithms.

Several researchers [Acar et al. 2000; Spoonhower et al. 2009] have shown that introducing parallelism through the use of futures can sometimes substantially reduce cache locality. In the worst case, if we add futures to a sequential program, a parallel execution managed by a work-stealing scheduler can incur $\Omega(PT_\infty + tT_\infty)$ deviations, which, as we show, implies $\Omega(CPT_\infty + CtT_\infty)$ more cache misses than the sequential execution. Here, C is the number of cache lines, P is the number of processors, t is the number of touches, and T_∞ is the computation's *span* (or *critical path*). As technology trends cause the cost of cache misses to increase, this additional cost is troubling.

This article makes the following three contributions. First, we show that if futures are used in a simple, disciplined way, then the situation with respect to cache locality is much better: if each future is touched only once, either by the thread that created it or by a later descendant of that thread, then parallel executions with work stealing can incur at most $O(CPT_\infty^2)$ additional cache misses, which is a substantial improvement over the unstructured case. This result provides a simple way to identify computations for which introducing futures will not incur a high cost in cache locality, as well as providing guidelines for the design of future-parallel computations. (Informally, we think these guidelines are natural and correspond to structures that programmers are likely to use anyway.) We also prove that this upper bound is tight within a factor of C . Our second contribution is to observe that when the scheduler has a choice between running the thread that created a future and the thread that implements the future, running the future thread first provides better cache locality. Finally, we show that certain variations of structured computation also have good cache locality.

The article is organized as follows. Section 2 describes the model for future-parallel computations. In Section 3, we describe parsimonious work-stealing schedulers and briefly discuss their cache performance measures. In Section 4, we define some restricted forms of structured future-parallel computations. Among them, we highlight structured single-touch computations, which, we believe, are likely to arise naturally in many programs. In Section 5.1, we prove that work-stealing schedulers on structured single-touch computations incur only $O(CPT_\infty^2)$ additional cache misses, if a processor always chooses the future to execute first when it creates that future. We also prove that this bound is tight within a factor of C . In Section 5.2, we show that if a processor chooses the current thread over the future thread when it creates that future, then the cache locality of a structured single-touch computation can be much worse. In Section 6, we show that some other kinds of structured future-parallel computations also achieve relatively good cache locality. Finally, we present conclusions in Section 7.

2. MODEL

In *fork-join parallelism* [Blelloch et al. 1995; Blelloch 1996; Blumofe et al. 1996], a sequential program is split into a directed acyclic graph (DAG) of *tasks* linked by directed dependency edges. These tasks are executed in an order consistent with their dependencies, and tasks unrelated by dependencies can be executed in parallel. Fork-join parallelism is well suited to dynamic load-balancing techniques such as *work stealing* [Burton and Sleep 1981; Halstead 1984, 1985; Arora et al. 1998; Blumofe and Leiserson 1999; Acar et al. 2000; Blumofe et al. 1995; Frigo et al. 1998; Kranz et al. 1989; Agrawal et al. 2007; Chase and Lev 2005].

A popular and effective way to extend fork-join parallelism is to allow threads to create *futures* [Halstead 1984, 1985; Arvind et al. 1989; Blelloch and Reid-Miller 1997; Fluet et al. 2008]. A future is a data object that represents a *promise* to deliver the result of an asynchronous computation when it is ready. That result becomes available to a thread when the thread *touches* that future, blocking if necessary until the result

is ready. Futures are attractive because they provide greater flexibility than fork-join programs, and they can also be implemented effectively using dynamic load-balancing techniques such as work stealing. Fork-join parallelism can be viewed as a special case of future parallelism, where the spawn operation is an implicit future creation and the sync operation is an implicit touch of the untouched futures created by a thread. Future parallelism is more flexible than fork-join parallelism because the programmer has finer-grained control over touches (joins).

2.1. Computation DAG

A thread creates a future by marking an expression (usually a method call) as a *future*. This statement spawns a new thread to evaluate that expression in parallel with the thread that created the future. When a thread needs access to the results of the computation, it applies a *touch* operation to the future. If the result is ready, it is returned by the touch, and otherwise the touching thread blocks until the result becomes ready. Without loss of generality, we will consider fork-join parallelism to be a special case of future parallelism, where forking a thread creates a future, and joining one thread to another is a touch operation.

Our notation and terminology follow earlier work [Arora et al. 1998; Blumofe and Leiserson 1999; Acar et al. 2000; Spoonhower et al. 2009]. A future-parallel computation is modeled as a DAG. Each node in the DAG represents a task (one or more instructions), and an edge from node u to node v represents the dependency constraint that u must be executed before v . We follow the convention that each node in the DAG has in-degree and out-degree of either 1 or 2, except for a distinguished *root node* with in-degree 0, where the computation starts, and a distinguished *final node* with out-degree 0, where the computation ends.

There are three types of edges:

- Continuation edges*, which point from one node to the next in the same thread,
- Future edges* (sometimes called *spawn edges*), which point from node u to the first node of another thread spawned at u by a future creation, and
- Touch edges* (sometimes called *join edges*), directed from a node u in one thread t to a node v in another thread, indicating that v touches the future computed by t .

A *thread* is a maximal chain of nodes connected by continuation edges. There is a distinguished *main thread* that begins at the root node and ends at the final node, and every other thread t begins at a node with an incoming future edge from a node of the thread that spawns t . The last node of t has only one outgoing edge, which is a touch edge directed to another thread, whereas other nodes of t may or may not have incoming and outgoing touch edges. A *critical path* of a DAG is a longest directed path in the DAG, and the DAG's *computation span* is the length of a critical path.

As illustrated in Figure 1, if a thread t_1 spawns a new thread t_2 at node v in t_1 (i.e., v has two outgoing edges, a continuation edge and a future edge to the first node of t_2), then we call t_1 the *parent thread* of t_2 , t_2 the *future thread* (of t_1) at v , and v the *fork* of t_2 . A thread t_3 is a *descendant thread* of t_1 if t_3 is a future thread of t_1 or, by induction, t_3 's parent thread is a descendant thread of t_1 .

If there is a touch edge directed from node v_1 in thread t_1 to node v_2 in thread t_2 (i.e., t_2 touches a future computed by t_1), and a continuation edge directed from node u_2 in t_2 to v_2 , then we call node v_2 a *touch of t_1 by t_2* , v_1 the *future parent* of v_2 , u_2 the *local parent* of v_2 , and t_1 the *future thread* of v_2 . (Note that the touch v_2 is actually a node in thread t_2 .) We call the fork of t_1 the *corresponding fork* of v_2 .

Note that only touch nodes have in-degree 2. To distinguish between the two types of nodes with out-degree 2—forks and future parents of touches—we follow the convention of previous work that the children of a fork both have in-degree 1 and cannot be touches.




Fig. 1. Node and thread terminology.

In this way, a fork node has two children with in-degree 1, whereas a touch's future parent has a (touch) child with in-degree 2.

We follow the convention that when a fork appears in a DAG, the future thread is shown on the left and the future parent on the right. (Note that this does not mean that the future thread is chosen to execute first at a fork.) Similarly, the future parent of a touch is shown on the left and the local parent on the right.

We use the following (standard) notation. Given a computation DAG, P is the number of processors executing the computation; t is the number of touches in the DAG; T_∞ , the *computation span* (or *critical path*), is the length of the longest directed path; and C is the number of cache lines in each processor.

3. WORK STEALING AND CACHE LOCALITY

In the article, we focus on parsimonious work-stealing algorithms [Arora et al. 1998], which have been extensively studied [Arora et al. 1998; Blumofe and Leiserson 1999; Acar et al. 2000; Spoonhower et al. 2009; Blumofe and Leiserson 1998] and used in systems such as Cilk [Blumofe et al. 1995]. In a parsimonious work-stealing algorithm, each processor is assigned a double-ended queue (deque). After a processor executes a node with out-degree 1, it continues to execute the next node if the next node is ready to execute. After the processor executes a fork, it pushes one child of the fork onto the bottom of its deque and executes the other. When the processor runs out of nodes to execute, it pops the first node from the bottom of its deque if the deque is not empty. If, however, its deque is empty, it steals a node from the top of the deque of an arbitrary processor.

In our model, a cache is fully associative and consists of multiple *cache lines*, each of which holds the data in a *memory block*. Each instruction can access only one memory block. In our analysis, we focus only on the widely used least-recently used (LRU) cache replacement policy, but our results about the upper bounds on cache overheads should apply to all *simple* cache replacement policies [Acar et al. 2000].¹

The *cache locality* of an execution is measured by the number of cache misses it incurs, which depends on the structure of the computation. To measure the effect on cache locality of parallelism, it is common to compare cache misses encountered in a sequential execution to the cache misses encountered in various parallel executions, focusing on the number of *additional* cache misses introduced by parallelism.

¹That is because the upper bounds in this article are based on the results of Acar et al. [2000] that bound the number of drifted nodes (i.e., deviations), and those results hold for all simple cache replacement policies, even with set associative caches, as discussed in Acar et al. [2000].




Fig. 2. The interesting part of the bound is $\Omega(CtT_\infty)$. Figure 5 in Spoonhower et al. [2009] shows a DAG, as a building block of a worst-case computation, that can incur $\Omega(T_\infty)$ deviations because of one touch. We can replace it with the DAG in Figure 2, which can incur $\Omega(CT_\infty)$ additional cache misses due to one touch v (if the processor at a fork always chooses the parent thread to execute first), so that the worst-case computation in Spoonhower et al. [2009] can incur $\Omega(CtT_\infty)$ additional cache misses because of t such touches. This DAG is similar to the DAG shown later in Figure 7(a) in this article. The proof of Theorem 5.7 shows how a parallel execution of this DAG incurs $\Omega(CT_\infty)$ additional cache misses.

Scheduling choices at forks affect the cache locality of executions with work stealing. After executing a fork, a processor picks one of the two child nodes to execute and pushes the other into its deque. For a sequential execution, whether a choice results in a better cache performance is a characteristic of the computation itself. For a parallel execution of a computation satisfying certain properties, however, we will show that choosing future threads (the left children) at forks to execute first guarantees a relatively good upper bound on the number of additional cache misses compared to a sequential execution that also chooses future threads first. In contrast, choosing the parent threads (the right children) to execute first can result in a large number of additional cache misses compared to a sequential execution that also chooses parent threads first.

4. STRUCTURED COMPUTATIONS

Consider a sequential execution where node v_1 is executed immediately before node v_2 . A *deviation* [Spoonhower et al. 2009], also called a *drifted node* [Acar et al. 2000], occurs in a parallel execution if a processor P executes v_2 , but not immediately after v_1 . For example, p might execute v_1 after v_2 , it might execute other nodes between v_1 and v_2 , or v_1 and v_2 might be executed by distinct processors.

Spoonhower et al. [2009] showed that a parallel execution of a future-parallel computation with work stealing can incur $\Omega(PT_\infty + tT_\infty)$ deviations. This implies that a parallel execution of a future-parallel computation with work stealing can incur $\Omega(PT_\infty + tT_\infty)$ additional cache misses. With minor modifications in that computation (Figure 2), a parallel execution can even incur $\Omega(CPT_\infty + CtT_\infty)$ additional cache misses.

Our contribution in this article is based on the observation that such poor cache locality occurs primarily when futures in the DAG can be touched by arbitrary threads, resulting in unrealistic and complicated dependencies. For example, in the worst-case




Fig. 3. A simplified version of the DAG in Spoonhower et al. [2009] that can incur high cache overhead. Here, v_1 and v_2 are touches. Suppose that a processor p_1 executes the root node, pushes the right child x of the root node into its deque, and then falls asleep. Now another processor p_2 steals x from p_1 's deque and executes the subgraph rooted at x . Thus, v_1 and v_2 will be checked (to see if they are available) even before the corresponding future threads are spawned at u_1 and u_2 .




Fig. 4. In this structured (single-touch) computation, the touches v_1 and v_2 will not be checked until their corresponding future threads have been spawned at u_1 and u_2 , respectively.

DAGs in Spoonhower et al. [2009] that can incur significantly high cache overheads, futures are touched by threads that can be created before the future threads computing these futures were created. As illustrated in Figure 3, a parallel execution of such a computation can arrive at a scenario where a thread touches a future before the future thread computing that future has been spawned. (As a practical matter, an implementation must ensure that such a touch does not return a reference to a memory location that has not yet been allocated.) Such scenarios are avoided by *structured* future-parallel computations (Figure 4) that follow certain simple restrictions.

Definition 4.1. A DAG is a *structured future-parallel computation* if (1) for the future thread t of any fork v , the local parents of the touches of t are descendants of v , and (2) at least one touch of t is a descendant of the right child of v .

There are two reasons we require that at least one touch of t is a descendant of the right child of v . First, it is natural that a computation spawns a future thread to compute a future because the computation itself later needs that value. At the fork v , the parent thread (the right child of v) represents the “main body” of the computation. Hence, the future will usually be touched either by the parent thread or by threads spawned directly or indirectly by the parent thread.

Second, a computation usually needs a kind of “barrier” synchronization to deal with resource release at the end of the computation. Some node in the future thread t , usually the last node, should have an outgoing edge pointing to the “main body” of the computation to tell the main body that the future thread has finished. Without such synchronization, t and its descendants will be isolated from the main body of the computation, and we can imagine a dangerous scenario where the main body of the computation finishes and releases its resources while t or its descendant threads are still running.

In our DAG model, such a synchronization point is by definition a touch node, although it may not be a real touch. We follow the convention that the thread that spawns

<pre> void MethodA { Future x = some computation; Future y = some computation; a = y.touch(); b = x.touch(); } </pre>	<pre> void MethodB { Future x = some computation; fork MethodC(x); } void MethodC(Future f) { a = f.touch(); } </pre>
(a)	(b)

Fig. 5. Two examples illustrating that single-touch computations are more flexible than fork-join computations.

a future thread releases it, so the synchronization point is a node in the parent thread or one of its descendants. Another possibility is to place the synchronization point at the last node of the entire computation, which is the typically case in languages such as Java, where the main thread of a program is in charge of releasing resources for the entire computation. These two styles are essentially equivalent and should have almost the same bounds on cache overheads. We will briefly discuss this issue in Section 6.2.

We consider how the following constraint affects cache locality.

Definition 4.2. A structured *single-touch* computation is a structured future-parallel computation, where each future thread spawned at a fork v is touched only once and the touch node is a descendant of v 's right child.

By the definition of threads, the future parent of the only touch of a future thread is the last node of that future thread (the last node can also be a parent of a join node, but we do not distinguish between a touch node and a join node). The DAG in Figure 4 represents a structured single-touch computation. We will show that work-stealing parallel executions of structured single-touch computations achieve significantly less cache overheads than unstructured computations.

In principle, a future could be touched multiple times by different threads, so structured single-touch computations are more restrictive structured computations in general. Nevertheless, the single-touch constraint is one that is likely to be observed by many programs. For example, as noted, the Cilk [Blumofe et al. 1995] language supports fork-join parallelism—a strict subset of the future-parallelism model considered here. If we interpret the Cilk language's `spawn` statement as creating a future, and its `sync` statement as touching all untouched futures previously created by that thread, then Cilk programs (like all fork-join programs) are structured single-touch computations.

Structured single-touch computations encompass fork-join computations but are strictly more flexible. Figure 5 presents two examples that illustrate the differences. If a thread creates multiple futures first and touches them later, fork-join parallelism requires they be touched (evaluated) in the reverse order. MethodA in Figure 5(a) shows the only order in which a thread can first create two futures and then touch them in a fork-join computation. This rules out, for instance, a program where a thread creates a sequence of futures, stores them in a priority queue, and evaluates them in some priority order. In contrast, our structured computations permit such futures to be evaluated by their creating thread or its descendants in any order.

Additionally, unlike fork-join parallelism, our notion of structured computation permits a thread to pass a future to a subroutine or descendant thread that touches that future, as illustrated in Figure 5(b). Our restrictions are that (1) only one thread can touch a future, and (2) the descendant thread that touches the future has to be created after the future. In fact, MethodC can even pass the future to a descendant of its own.

In a fork-join computation, however, only the thread creating the future can touch it, which is much more restrictive. We believe these restrictions are easy to follow and should be compatible with how many people program in practice.

Blelloch and Reid-Miller [1997] observe that if a future can be touched multiple times, then complex and potentially inefficient operations and data structures are needed to correctly resume the suspended threads that are waiting for the touch. By contrast, the runtime support for futures can be significantly simplified if each future is touched at most once.

We also consider the following structured local-touch computations in the article.

Definition 4.3. A structured *local-touch* computation is one where each future thread spawned at a fork v is touched only at nodes in its parent thread, and these touches are descendants of the right child of v .

Informally, the local touch constraint implies that a thread that needs the value of a future should create the future itself. Note that in a structured computation with local-touch constraint, a future thread is now allowed to evaluate multiple futures and these futures can be touched at different times. Although allowing a future thread to compute multiple futures is not very common, Blelloch and Reid-Miller [1997] point out that it can be useful for some future-parallel computations like pipeline parallelism [Blelloch and Reid-Miller 1997; Blumofe et al. 1995; Gordon et al. 2006; Giacomoni et al. 2008; Lee et al. 2013]. We will show in Section 6.1 that work-stealing parallel executions of computations satisfying the local-touch constraint also have relatively low cache overheads. Note that structured computations with both single-touch and local-touch constraints are still a superset of fork-join computations.

5. STRUCTURED SINGLE-TOUCH COMPUTATIONS

5.1. Future Thread First at Each Fork

We now analyze cache performance of work stealing on parallel executions of structured single-touch computations. We will show that work stealing has relatively low cache overhead if the processor at a fork always chooses the future thread to execute first and puts the parent future into its deque. For brevity, all arguments and results in this section assume that every execution chooses the future thread at a fork to execute first.

LEMMA 5.1. *In the sequential execution of a structured single-touch computation, any touch x 's future parent is executed before x 's local parent, and the right child of x 's corresponding fork v immediately follows x 's future parent.*

PROOF. By induction. Given a DAG, initially let S be an empty set and T the set of all touches. Note that

$$S \cap T = \emptyset \text{ and } S \cup T = \{\text{all touches}\}. \quad (1)$$

Consider any touch x in T such that x has no ancestors in T . (In other words, x has no ancestor nodes that are also touches.) Let t be the future thread of x and v the corresponding fork. Note that x 's future parent is the last node of t by definition. When the single processor executes v , the processor pushes v 's right child into the deque and continues to execute thread t . By hypothesis, there are no touches by t , as any touch by t must be an ancestor of x . There may be some forks in t . However, whenever the single processor executes a fork in t , it pushes the right child of that fork, which is a node in t , into the deque, and hence t (i.e., a node in t) is right below v 's right child in the deque. Therefore, the processor will always resume thread t before the right child of v . Since there is no touch by t , all nodes in t are ready to execute one by one. Thus,

when the future parent of the touch x is executed eventually, the right child of v is right at the bottom of the deque. By the single-touch constraint, the local parent of x is a descendant of the right child of v , so the local parent of x cannot be executed yet. Thus, the processor will now pop the right child of v out from the bottom of the deque. Since this node is not a touch, it is ready to execute. Therefore, x satisfies the following two properties.

PROPERTY 5.2. *Its future parent is executed before its local parent.*

PROPERTY 5.3. *The right child of its corresponding fork immediately follows its future parent.*

Now set $S = S \cup \{x\}$ and $T = T - \{x\}$. Thus, all touches in S satisfy Properties 5.2 and 5.3. Note that Equation (1) still holds.

Now suppose that at some point all nodes in S satisfy Properties 5.2 and 5.3, and that Equation (1) holds. Again, we now consider a touch x in T such that no touches in T are ancestors of x —that is, all touches that are ancestors of x are in S . Since the computation graph is a DAG, there must be such an x as long as T is not empty. Let t be the future thread of x and v the corresponding fork. If there are no touches by t , then we can prove that x satisfies Properties 5.2 and 5.3 by the same argument for the first touch added into S . Now assume that there are touches by t . Since those touches are ancestors of x , they are all in S and hence they all satisfy Property 5.2. When the processor executes v , it pushes v 's right child into the deque and starts executing t . Similar to what we showed earlier, when the processor gets to a fork in t , it will always push t into its deque, right below the right child of v . Thus, the processor will always resume t before the right child of v . When the processor gets to the local parent of a touch by t , we know that the future parent of the touch has already been executed since the touch satisfies Property 5.2. Thus, the processor can immediately execute that touch and continue to execute t . Therefore, the processor will eventually execute the future parent of x , whereas the right child of t is still the next node to pop in the deque. Again, since the local parent of x is a descendant of the right child of v , the local parent of x as well as x cannot be executed yet. Therefore, the processor will now pop the right child of v to execute, and hence x satisfies Properties 5.2 and 5.3. Now we set $S = S \cup \{x\}$ and $T = T - \{x\}$. Therefore, all touches in S satisfy Properties 5.2 and 5.3, and Equation (1) also holds. By induction, we have $S = \{\text{all touches}\}$, and all touches satisfy Properties 5.2 and 5.3. \square

Acar et al. [2000] have shown that the number of additional cache misses in a work-stealing parallel computation is bounded by the product of the number of deviations and the number of cache lines. It is easy to see that only two types of nodes in a DAG can be deviations: the touches and the child nodes of forks that are not chosen to execute first. Since we assume the future thread (left child) at a fork is always executed first, only the right children of forks can be deviations. Next, we bound the number of deviations incurred by a work-stealing parallel execution to bound its cache overhead.

LEMMA 5.4. *Let t be the future thread at a fork v in a structured single-touch computation. If t 's touch x or v 's right child u is a deviation, then either u is stolen or there is a touch by t which is a deviation.*

PROOF. By Lemma 5.1, a touch is a deviation if and only if its local parent is executed before its future parent. Now suppose that a processor p executes v and pushes u into its deque. Assume that u is not stolen and that no touches by t are deviations. Thus, u will stay in p 's deque until p pops it out. The proof of this lemma is similar to that of Lemma 5.1. After p spawns thread t at v , it moves to execute t . When p executes

“ordinary” nodes in t , no nodes are pushed into or popped out of p ’s deque, and hence u is still the next node in the deque to pop. When p executes a fork in t , it pushes t (more specifically, the right child of that fork) into its deque, right below u . Since a thief processor always steals from the top of a deque, and by hypothesis u is not stolen, t cannot be stolen. Thus, p will always resume t before u and then u will become the next node in the deque to pop. When p executes the local parent of a touch by t , the future parent of that touch must have been executed, as we assume that touch is not a deviation. Thus, p can continue to execute that touch immediately and keep moving on in t with its deque unchanged. Therefore, p will finally get to the local parent of x and then pop u out from its deque, as x is a descendant of u and x cannot be executed yet. Hence, neither x nor u can be a deviation. \square

THEOREM 5.5. *If, at each fork, the future thread is chosen to execute first, then a parallel execution with work stealing incurs $O(PT_\infty^2)$ deviations and $O(CPT_\infty^2)$ additional cache misses in expectation on a structured single-touch computation, where (as usual) P is the number of processors involved in this computation, T_∞ is the computation span, and C is the number of cache lines.*

PROOF. Arora et al. [1998] have shown that in a parallel execution with work stealing, there are in expectation $O(PT_\infty)$ steals. Now let us count how many deviations these steals can incur. A steal on the right child u of a fork v can make u and v ’s corresponding touch x_1 deviations. Suppose that x_1 is a touch by a thread t_2 , then the right child of the fork of t_2 and t_2 ’s touch x_2 can be deviations. If x_2 is a deviation and x_2 is a touch by another thread t_3 , then the right child of the fork of t_3 and t_3 ’s touch x_3 can be deviation as well. Note that x_2 is a descendant of x_1 , and x_3 is a descendant of x_2 . By repeating this observation, we can find a chain of touches $x_1, x_2, x_3, \dots, x_n$, called a *deviation chain*, such that each x_i and the right child of the corresponding fork of x_i can be deviations. Since for each $i > 1$, x_i is a descendant of $x_2, x_1, x_2, x_3, \dots, x_n$ is in a directed path in the computation DAG. Since the length of any path is at most T_∞ , we have $n \leq T_\infty$. Since each future thread has only one touch, there is only one deviation chain for a steal. Since there are $O(PT_\infty)$ steals in expectation in a parallel execution [Arora et al. 1998], we can find in expectation $O(PT_\infty)$ deviation chains and in total $O(PT_\infty^2)$ touches and right children of the corresponding forks involved (i.e., $O(PT_\infty^2)$ deviations involved).

Next, we prove by contradiction that no other touches or right children of forks can be deviations. Suppose that there is touch y such that y or the right child of the corresponding fork of y is a deviation and that y is not in any deviation chain. The right child of the corresponding fork of y cannot be stolen, since by hypothesis y is not the first touch in any of those chains. Thus, by Lemma 5.4, there is a touch y' by the future thread of y and y' is a deviation. Note that y' ’s cannot be in any deviation chain either. Otherwise, y and the deviation chain that y' is in will form a deviation chain as well, a contradiction. Therefore, by repeating such “tracing back,” we will end up at a deviation touch that is not in any deviation chain and has no touches as its ancestors. Therefore, there are no touches by the future thread of this touch, and the right child of the corresponding future fork of it is not stolen, contradicting Lemma 5.4.

The upper bound on the expected number of additional cache misses follows from the result of Acar et al. [2000] that the number of additional cache misses in a work-stealing parallel computation is bounded by the product of the number of deviations and the number of cache lines. \square

The bound on the number of deviations in Theorem 5.5 is tight, and the bound on the number of additional cache misses is tight within a factor of C , as shown next in Theorem 5.6.




Fig. 6. Part (c) shows a DAG on which work stealing can incur $\Omega(PT_\infty^2)$ deviations and $\Omega(PT_\infty^2)$ additional cache misses. It uses the DAGs in (a) and (b) as building blocks.

THEOREM 5.6. *If, at each fork node, the future thread is chosen to execute first, then a parallel execution with work stealing can incur $\Omega(PT_\infty^2)$ deviations and $\Omega(PT_\infty^2)$ additional cache misses on a structured single-touch computation, whereas the sequential execution of this computation incurs $O(PT_\infty^2/C)$ cache misses.*

PROOF. Figure 6(c) shows a computation DAG on which we can get the bounds that we want to prove. The DAG in Figure 6(c) uses the DAGs in Figure 6(a) and (b) as building blocks. Let us look at Figure 6(a) first. Suppose that there are two processors p_1 and p_2 executing the DAG in Figure 6(a). Suppose that p_2 executes v , pushes u_1 into its deque, and then falls asleep before executing w . Now suppose that p_1 steals u_1 . For each $i \leq k$, neither s_i nor Z_i can be executed, as w has not been executed yet. Now p_1 takes a solo run, executing $u_1, x_1, Y_1, u_2, x_2, Y_2, \dots, x_k, Y_k$. After p_1 finishes, p_2 wakes up and executes the rest of the computation DAG. Note that the right (local) parent of s_i is executed before the left (future) parent of the touch is executed. Thus, by Lemma 5.1, each s_i is a deviation. Hence, this parallel execution incurs k deviations, and the computation span of the computation is $\Theta(k)$.

Now let us consider a parallel execution of the computation in Figure 6(b). For each $i \leq k$, the subgraph rooted at v_i is identical to the computation DAG in Figure 6(a) (except the last node of the subgraph has an extra edge pointing to a node of the main thread). Suppose that there are three processors, p_1 , p_2 , and p_3 , working on the computation. Assume that p_2 executes r_1 and v_1 and then falls asleep when it is about to execute w . p_3 now steals r_2 from p_2 and then falls asleep as well. Then p_1 steals u_1 from p_2 's deque. Now p_1 and p_2 execute the subgraph rooted at v_1 in the same way they execute the DAG in Figure 6(a). After p_1 and p_2 finish, p_3 wakes up and executes r_2 . Now these three processors start working on the subgraph rooted at r_3 in the same way they executed the graph rooted at r_1 . By repeating this, the execution ends up incurring k^2 deviations when all k subgraphs are done. Since the length of the path $r_1, r_2, r_3 \dots$ on the right-hand side is $\Theta(k)$, the computation span of the DAG is still $\Theta(k)$.

Now we construct the final computation DAG, as in Figure 6(c). The “top” nodes of the DAG are all forks, each spawning a future thread. Thus, they form a binary tree, and the number of threads increase exponentially. The DAG stops creating new threads at

level $\Theta(\log n)$ when it has n threads rooted at S_1, S_2, \dots, S_n , respectively. For each i , the subgraph rooted at S_i is identical to the DAG in Figure 6(b). Suppose that there are $3n$ processors working on the computation. It is easy to see that n processors can eventually get to S_1, S_2, \dots, S_n . Suppose that they all fall asleep immediately after executing the first two nodes of S_i (corresponding to r_1 and v_1 in Figure 6(b)) and then each two of the rest of the $2n$ free processors join to work on the subgraph rooted at S_i , in the same way p_1, p_2 , and p_3 did in Figure 6(b). Therefore, this execution will finally incur nk^2 deviations, whereas the computation span of the DAG is $\Theta(k + \log n)$. Therefore, by setting $n = P/3$, we get a parallel execution that incurs $\Omega(PT_\infty^2)$ deviations when $\log P = O(k)$.

To get the bound on the number of additional cache misses, we just need to modify the graph in Figure 6(a) as follows. For each $1 \leq i \leq k$, Y_i consists of a chain of C nodes $y_{i1}, y_{i2}, \dots, y_{iC}$, where C is the number of cache lines. $y_{i1}, y_{i2}, \dots, y_{iC}$ access memory blocks m_1, m_2, \dots, m_C , respectively. Similarly, each Z_i consists of a chain of C nodes $z_{i1}, z_{i2}, \dots, z_{iC}$. $z_{i1}, z_{i2}, \dots, z_{iC}$ access memory blocks m_C, m_{C-1}, \dots, m_1 , respectively. All s_i access memory block m_C . For all $1 \leq i \leq k$, u_i and x_i both access memory block m_{C+1} . It does not matter which memory blocks the other nodes in the DAG access. For simplicity, assume that the other nodes do not access memory. In the sequential execution, the single processor has m_1, m_2, \dots, m_C in its cache after executing v, w, u_1, x_1, Y_1, Z_1 and has incurred $(C + 1)$ cache misses so far. Now it executes u_2 and x_2 , incurring one cache miss at node u_2 by replacing m_C with m_{C+1} in its cache, as m_C is the LRU block. When it executes Y_2 and Z_2 , it only incurs one cache miss by replacing m_{C+1} with m_C at the last node of Y_2, y_{2C} . Likewise, it is easy to see that the sequential execution will only incur cache misses at nodes u_i and at the last nodes of Y_i for all i . Hence, the sequential execution incurs only $O(k + C)$ cache misses. When $k = \Omega(C)$, the sequential execution incurs only $O(k)$ cache misses.

Now consider the parallel execution by two processors p_1 and p_2 that we described before. p_2 will incur only C cache misses, as Z_i and s_i only access m different blocks m_1, m_2, \dots, m_C and hence p_2 does not need to swap any memory blocks out of its cache. However, p_1 will incur lots of cache misses. After executing each Y_i , p_1 will execute u_{i+1} . Thus, at u_{i+1} , one cache miss is incurred and m_1 is replaced with m_{C+1} , as m_1 is the LRU block. Then when p_1 executes the first node $y_{(i+1)1}$ in Y_i , m_1 is not in its cache. Since m_2 now becomes the LRU memory block in p_1 's cache, m_2 is replaced by m_1 . Thus, m_2 will not be in the cache when it is in need at $y_{(i+1)2}$. Therefore, it is obvious that p_1 will incur a cache miss at each node in Y_i and hence incur Ck cache misses in total in the entire execution. Note that the computation span of this modified DAG is $\Theta(Ck)$, as each Z_i now has C nodes. Therefore, the sequential execution and the parallel execution actually incur $\Theta(T_\infty/C)$ and $\Theta(T_\infty)$, respectively, when $\log P = O(k)$. Therefore, if we use this modified DAG as the building blocks in Figure 6(c), we will get the bound on the number of additional cache misses stated in the theorem. \square

5.2. Parent Thread First at Each Fork

In this section, we show that if the parent thread is always executed first at a fork, a work-stealing parallel execution of a structured single-touch computation can incur $\Omega(tT_\infty)$ deviations and $\Omega(CtT_\infty)$ additional cache misses, where t is the number of touches in the computation, whereas the corresponding sequential execution incurs only a small number of cache misses. This bound matches the upper bound for general, unstructured future-parallel computations [Spoonhower et al. 2009].² This result,

²The bound on the expected number of deviations in Spoonhower et al. [2009] is actually $O(PT_\infty + tT_\infty)$. However, as pointed out in Spoonhower et al. [2009], a simple fork-join computation can get $\Omega(PT_\infty)$ deviations. Hence, we focus on the more interesting part $\Omega(tT_\infty)$.




Fig. 7. DAGs used by Figure 8 as building blocks.

combined with the result in Section 5.1, shows that choosing the future threads at forks to execute first achieves better cache locality for work-stealing schedulers on structured single-touch computations.

THEOREM 5.7. *If, at each fork, the parent thread is chosen to execute first, then a parallel execution with work stealing can incur $\Omega(tT_\infty)$ deviations and $\Omega(CtT_\infty)$ additional cache misses on a structured single-touch computation, whereas the sequential execution of this computation incurs only $O(C + t)$ cache misses.*

PROOF. The final DAG that we want to construct is shown later in Figure 8. It uses the DAGs in Figure 7 as building blocks. We first describe how a single deviation at a touch u_3 can incur $\Omega(T_\infty)$ deviations and $\Omega(CT_\infty)$ additional cache misses in Figure 7(a). To get the bound that we want to prove, here we follow the convention in Acar et al. [2000] and Spoonhower et al. [2009] to distinguish between touches and join nodes in the DAG. More specifically, y_i is a join node, not a touch, for each $1 \leq i \leq n$. For each

$1 \leq i \leq n$, node x_i accesses memory block m_1 and y_i accesses memory block m_{C+1} . Z_i consists of a chain of C nodes $z_{i1}, z_{i2}, \dots, z_{iC}$, accessing memory blocks m_1, m_2, \dots, m_C , respectively. All other nodes do not access memory. Assume that in the sequential execution a single processor p_1 executes the entire DAG in Figure 7(a). Suppose that initially the left (future) parent of u_3 has already been executed. p_1 starts executing the DAG at u_1 . Since p_1 always stays on the parent thread at a fork, it first pushes s into its deque, continues to execute u_2, u_3, u_4 , and then executes x_1, x_2, \dots, x_n while pushing $z_{11}, z_{21}, \dots, z_{n1}$ into its deque. Since v cannot be executed due to s , p_1 pops z_{n1} out of its deque and executes the nodes in Z_n . Then p_1 executes all nodes in $Z_{n-1}, Z_{n-2}, \dots, Z_1$, in this order. So far, p_1 has only incurred C cache misses, as all nodes that it has executed only access memory blocks m_1, \dots, m_C , and hence it did not need to swap any memory blocks out of its cache. Now p_1 executes s, v and then y_n, y_{n-1}, \dots, y_1 , incurring only one more cache miss by replacing m_1 with m_{C+1} at y_n . Hence, this execution incurs $O(C)$ cache misses in total. Note that the left parent of y_i is executed before the right parent y_i for all i .

Now assume that in another execution by p_1 , the left parent of u_3 is in p_1 's deque when p_1 starts executing u_1 . Thus, u_3 is a deviation with respect to the previous execution. Since u_3 is not ready to execute after p_1 executes u_2 , p_1 pops s out of its deque to execute. Since v is not ready, p_1 now pops the left parent of u_3 to execute and then executes $u_3, u_4, x_1, x_2, \dots, x_n, v$. Now p_1 pops z_{n1} out and executes all nodes Z_n . Note that y_n is now ready to execute, and the memory blocks in p_1 's cache at the moment are m_1, m_2, \dots, m_C . Now p_1 executes y_n , replacing the LRU block m_1 with m_{C+1} . p_1 then pops $z_{(n-1)1}$ out and executes all nodes $z_{(n-1)1}, z_{(n-1)2}, \dots, z_{(n-1)C}$ in Z_{n-1} one by one. When p_1 executes $z_{(n-1)1}$, it replaces m_2 with m_1 , and when it executes $z_{(n-1)2}$, it replaces m_3 with m_2 , and so on. The same thing happens to all Z_i and y_i . Thus, p_1 will incur a cache miss at every node afterward, ending up with $\Omega(Cn)$ cache misses in total. Note that the computation span of this DAG is $T_\infty = \Theta(C + n)$. Thus, this execution with a deviation at u_3 incurs $\Omega(CT_\infty)$ cache misses when $n = \Omega(C)$. Moreover, all y_i are deviations, and hence this execution incurs $\Omega(T_\infty)$ deviations.

Now let us see how a single steal at the beginning of a thread results in $\Omega(T_\infty)$ deviations and $\Omega(CT_\infty)$ cache misses at the end of the thread. Figure 7(b) presents such a computation. First we consider the sequential execution by a processor p_1 . It is easy to see that p_1 executes nodes in the order $r, u_1, w_1, s_2, s_1, v_1, u_2, w_2, v_2, u_3, w_3, s_4, s_3, v_3, u_4, \dots$. The key observation is that w_i is executed before s_i is executed for any odd-numbered i , whereas w_i is executed after s_i is executed for any even-numbered i . This statement can be proved by induction. Obviously, this holds for $i = 1$ and $i = 2$, as we showed before. Now suppose that this fact holds for all $1, 2, \dots, i$, for some even-numbered i . Now suppose that p_1 executes u_{i-1} . Then p_1 pushes s_i into its deque and executes w_{i-1} . Since we know that w_{i-1} should be executed before s_{i-1} , s_{i-1} has not been executed yet. Moreover, s_{i-1} must already be in the deque before s_i was pushed into the deque, as s_{i-1} 's parent u_{i-2} has been executed and s_{i-1} is ready to execute. Now p_1 pops s_i out to execute. Since v_i is not ready to execute, p_1 pops s_{i-1} out and then executes s_{i-1}, v_{i-1}, u_i , and pushes s_{i+1} into the deque. Now p_1 continues to execute w_i, v_i, u_{i+1} and pushes s_{i+1} into its deque. Then p_i executes w_{i+1} and pops s_{i+2} out, as v_{i+1} is not ready due to s_{i+1} . Now we can see that w_{i+1} and s_{i+2} have been executed, but s_{i+1} and w_{i+2} have not yet been executed. In other words, w_{i+1} is executed before s_{i+1} and w_{i+2} is executed after s_{i+2} . Therefore, the statement holds for $i + 1$ and $i + 2$, and hence the proof completes.

The subgraph rooted at u_k is identical to the graph in Figure 7(a), with v_k corresponding to u_3 in Figure 7(a). Therefore, if k is an even number, v_k 's left parent has been executed when w_k is executed, and hence the sequential execution will incur only $O(C)$ cache misses on the subgraph rooted at u_k .




Fig. 8. A DAG on which work stealing can incur $\Omega(tT_\infty)$ deviations and $\Omega(CtT_\infty)$ if it chooses parents threads to execute first at forks. This example uses the DAGs in Figure 7 as building blocks.

Now consider the following parallel execution of the DAG in Figure 7(b) by two processors p_1 and p_2 . p_1 first executes r and pushes s_1 into its deque. Then p_2 immediately steals s_1 and executes it. Now p_2 falls asleep, leaving p_1 executing the rest of the DAG alone. It is easy to see that p_1 will execute the nodes in the DAG in the order $u_1, w_1, v_1, u_2, w_2, s_3, s_2, v_2, u_3, w_3, v_3, u_4, s_4, \dots$. It can be proved by induction that w_i is executed after s_i is executed for any odd-numbered i , whereas w_i is executed before s_i is executed for any even-numbered i , which is opposite to the order in the sequential execution. The induction proof is similar to that of the previous observation in the sequential execution, so we omit the proof here. If k is an even number, w_k will be executed before the left parent of v_k , and hence this execution will incur $\Omega(T_\infty)$ deviations and $\Omega(CtT_\infty)$ cache misses when $n = \Omega(C)$ and $n = \Omega(k)$.

The final DAG we want to construct is in Figure 8. This is actually a generalization of the DAG in Figure 7(b). Instead of having one fork u_i before each touch v_i , it has two forks u_i and x_i for each i . After each touch v_i , the thread at y_i splits into two identical branches, touching the futures spawned at u_i and x_i , respectively. In this figure, we only depict the right branch and omit the identical left branch. As we can see, the right branch later has a touch v_{i+1} touching the future s_{i+1} spawned at the fork x_i . If we only look at the thread on the right-hand side, it is essentially the same as the DAG in Figure 7(b). The sequential execution of this DAG by p_1 is similar to that in Figure 7(b). The only difference is that p_1 at each y_i will execute the right branch first and then the left branch recursively. Similarly, it can be proved by induction that w_i is executed before s_i is executed for any odd-numbered i , whereas w_i is executed after s_i is executed

for any even-numbered i . Obviously, this also holds for each left branch. Now consider a parallel execution by two processors p_1 and p_2 . p_1 first executes r . p_2 immediately steals s_1 and executes it and then sleeps forever. Now p_1 makes a solo run to execute the rest of the DAG. Again, we can prove by the same induction argument that w_i is executed after s_i is executed for any odd-numbered i , whereas w_i is executed before s_i is executed for any even-numbered i , which is opposite to the order in the sequential execution. The preceding two induction proofs are a little more complicated than those for the DAG in Figure 7(b), but the ideas are essentially the same (the only difference is now we have to prove that the statements hold for the two identical branches split at fork y_i at the inductive step), and hence we omit the proofs again.

By splitting each thread into two after each y_i , the number of branches in the DAG increases exponentially. Suppose that there are t touches in the DAG. Thus, there are eventually $\Theta(t)$ branches, and the height of this structure is $\Theta(\log t)$. At the end of each branch is a subgraph identical to the DAG in Figure 7(a). Therefore, the parallel execution with only one steal can end up incurring $\Theta(tn)$ deviations and $\Theta(Ctn)$ cache misses. The sequential execution incurs only $\Theta(C + t)$ cache misses, as the sequential execution will incur only two cache misses by swapping m_{C+1} in and out at each branch, after it incurs C cache misses to load m_1, m_2, \dots, m_C at the first branch. Hence, when $n = \Omega(\log t)$ and $n = \Omega(C)$, we get the bound stated in the theorem. \square

6. OTHER KINDS OF STRUCTURED COMPUTATIONS

It is natural to ask whether other kinds of structured computations can also achieve relatively good cache locality. We now consider two alternative kinds of restrictions.

6.1. Structured Local-Touch Computations

In this section, we prove that work-stealing parallel executions of structured local-touch computations also have relatively good cache locality, if the future thread is chosen to execute first at each fork. This result, combined with Theorems 5.5 and 5.7, implies that work-stealing schedulers for structured computations are likely better off choosing future threads to execute first at forks.

LEMMA 6.1. *In the sequential execution of a structured local-touch computation where the future thread at a fork is always chosen to execute first, any touch x 's future parent is executed before x 's local parent, and the right child of any fork v immediately follows the last node of the future thread spawned at v —that is, the future parent of the last touch of the future thread.*

The proof is omitted because it is almost identical to that of Lemma 5.1. (We first consider a future thread whose touches are the “earliest” in the DAG; in other words, no other touches are ancestors of them, and we can easily prove the statement in Lemma 6.1 that holds for those touches. Then by the same induction proof as for Lemma 5.1, we can prove that the statement holds for all future threads' touches.)

THEOREM 6.2. *If the future thread at a fork is always chosen to execute first, then a parallel execution with work stealing incurs $O(PT_\infty^2)$ deviations and $O(CPT_\infty^2)$ additional cache misses in expectation on a structured local-touch computation.*

PROOF. Let v be a fork that spawns a future thread t . Now we consider a parallel execution. Let p be a processor that executes v and pushes the right child of v into its deque. Suppose that the right child of v is not stolen. Now consider the subgraph G' consisting of t and its descendant threads. Note that G' itself is a structured computation DAG with local-touch constraint. Now p starts executing G' .

According to local-touch constraint, the only nodes outside G' that connect to the nodes in G' are v and the touches of t , and c is the only node outside G' that the nodes on which G' depend. Now v has been executed, and the touches of t are not ready to execute due to the right child of v . Hence, p is able to make a sequential execution on G' without waiting for any node outside to be done or jumping to a node outside, as long as no one steals a node in G' from p 's deque. Since we assume that the right child of v will not be stolen and any nodes in G' can only be pushed into p 's deque below v , no nodes in G' can be stolen. Hence, G' will be executed by a sequential execution by p . Therefore, there are no deviations in G' . After p executed the last node in G' , which is the last node in t , p pops the right child of v to execute. Hence, the right child of v cannot be a deviation either, if it is not stolen. In other words, those nodes can be deviations only if the right child of v is stolen. Since there are in expectation $O(PT_\infty)$ steals in an parallel execution and each future thread has at most T_∞ touches, the expected number of deviations is bounded by $O(PT_\infty^2)$ and the expected number of additional touches is bounded by $O(CPT_\infty^2)$. \square

6.2. Structured Computations with Super Final Nodes

As discussed in Section 4, in languages such as Java, the program's main thread typically releases all resources at the end of an execution. To model this structure, we add an edge from the last node of each thread to the final node of the computation DAG. Thus, the final node becomes the only node with in-degree greater than 2. Since the final node is always the last to execute, simply adding those edges pointing to the final node into a DAG will not change the execution order of the nodes in the DAG. It is easy to see that having such a super node will not change the upper bound on the cache overheads of the work-stealing parallel executions of a structured computation.

For structured computations with super final nodes, it also makes sense to slightly relax the single-touch constraint as follows.

Definition 6.3. A structured single-touch computation with a *super final node* is one where each future thread t at a fork v has at least one and at most *two* touches, a descendant of v 's right child and the super final node.

In such a computation, a future thread can have the super final node as its only touch. This structure corresponds to a program where one thread forks another thread to accomplish a side effect instead of computing a value. The parent thread never touches the resulting future, but the computation as a whole cannot terminate until the forked thread completes its work.

Now we show that the parallel executions of structured single-touch computations with super final nodes also have relatively low cache overheads.

LEMMA 6.4. *In the sequential execution of a structured single-touch computation with a super final node, where the future thread at a fork is always chosen to execute first, any touch x 's future parent is executed before x 's local parent, and the right child u of any fork v immediately follows the last node of the future thread spawned at v —that is, the future parent of the last touch of the future thread.*

LEMMA 6.5. *Let t be the future thread at a fork v in a structured single-touch computation with a super final node. If a touch of t or v 's right child u is a deviation, then either u is stolen or there is a touch by t , which is a deviation.*

PROOF. The proofs of Lemma 5.1 and Lemma 5.4, with only minor modifications, also apply to the preceding two lemmas, respectively, because introducing the super final node into a computation does not affect the order in which other nodes are executed, as no other nodes need to wait for the super final node and the super final node is always

the last node to execute. More specifically, when a processor executing any thread t reaches a node that is a parent of the super final node, the processor will continue to work on t if that node is not the last node of t , and otherwise try popping a node out of its deque. Therefore, by the same proof techniques as for Lemmas 5.1 and 5.4, we can show that a processor will execute the right child u of a fork v and the parents of the touches of the future spawned at v in the order stated in Lemmas 6.4 and 6.5. \square

THEOREM 6.6. *If, at each fork, the future thread is chosen to execute first, then a parallel execution with work stealing incurs $O(PT_\infty^2)$ deviations and $O(CPT_\infty^2)$ additional cache misses in expectation on a structured single-touch computation with a super final node.*

PROOF. The proof is similar to that of Theorem 5.5. The only difference is that if a touch by a thread t is a deviation, now the two touches of t can both be deviations, which could be a trouble for constructing the deviation chains. Fortunately, one of these two touches is the super final node, which is always the last node to execute and hence will not make the touches of other threads become deviations. Therefore, we can still get a unique deviation chain starting from a steal, and hence the proof of Theorem 5.5 still applies here. \square

Similarly, we can also introduce a super final node to a structured local-touch computation as follows.

Definition 6.7. A structured local-touch computation with a *super final node* is one where each future thread t spawned at a fork v can be touched only by the super final node and by t 's parent thread at nodes that are descendants of the right child of v .

It is obvious that by the same proof as for Theorem 6.2, we can prove the following bounds.

THEOREM 6.8. *If the future thread at a fork is always chosen to execute first, then a parallel execution with work stealing incurs $O(PT_\infty^2)$ deviations and $O(CPT_\infty^2)$ additional cache misses in expectation on a structured local-touch computation with a super final node.*

7. CONCLUSIONS

We have focused primarily on structured single-touch computations, in which futures are used in a restricted way. We saw that for such computations, a parallel execution by a work-stealing scheduler that runs future threads first can incur at most $O(CPT_\infty^2)$ cache misses more than the corresponding sequential execution—a substantially better cache locality than the $\Omega(CPT_\infty + CtT_\infty)$ worst-case additional cache misses possible with unstructured use of futures. Although we cannot prove this claim formally, we think that these restrictions correspond to program structures that would occur naturally anyway in many (but not all) parallel programs that use futures. For example, Cilk [Blumofe et al. 1995] programs are structured single-touch computations, and Blelloch and Reid-Miller [1997] observe that the single-touch requirement substantially simplifies implementations.

We also considered some alternative restrictions on future use, such as structured local-touch computations, and structured computations with super final nodes, which also incur a relatively low cache-locality penalty. In terms of future work, we think that it would be promising to investigate how far these restrictions can be weakened or modified while still avoiding a high cache-locality penalty. We would also like to understand how these observations can be exploited by future compilers and runtime systems.

REFERENCES

- Umut A. Acar, Guy E. Blelloch, and Robert D. Blumofe. 2000. The data locality of work stealing. In *Proceedings of the 12th Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA'00)*. ACM, New York, NY, 1–12. DOI: <http://dx.doi.org/10.1145/341800.341801>
- Kunal Agrawal, Yuxiong He, and Charles E. Leiserson. 2007. Adaptive work stealing with parallelism feedback. In *Proceedings of the 12th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP'07)*. ACM, New York, NY, 112–120. DOI: <http://dx.doi.org/10.1145/1229428.1229448>
- Nimar S. Arora, Robert D. Blumofe, and C. Greg Plaxton. 1998. Thread scheduling for multiprogrammed multiprocessors. In *Proceedings of the 10th Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA'98)*. ACM, New York, NY, 119–129. DOI: <http://dx.doi.org/10.1145/277651.277678>
- Arvind, Rishiyur S. Nikhil, and Keshav K. Pingali. 1989. I-structures: Data structures for parallel computing. *ACM Transactions on Programming Languages and Systems* 11, 4, 598–632. DOI: <http://dx.doi.org/10.1145/69558.69562>
- Guy E. Blelloch. 1996. Programming parallel algorithms. *Communications of the ACM* 39, 3, 85–97. DOI: <http://dx.doi.org/10.1145/227234.227246>
- Guy E. Blelloch, Phillip B. Gibbons, and Yossi Matias. 1995. Provably efficient scheduling for languages with fine-grained parallelism. In *Proceedings of the 7th Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA'95)*. ACM, New York, NY, 1–12. DOI: <http://dx.doi.org/10.1145/215399.215403>
- Guy E. Blelloch and Margaret Reid-Miller. 1997. Pipelining with futures. In *Proceedings of the 9th Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA'97)*. ACM, New York, NY, 249–259. DOI: <http://dx.doi.org/10.1145/258492.258517>
- Robert D. Blumofe, Matteo Frigo, Christopher F. Joerg, Charles E. Leiserson, and Keith H. Randall. 1996. An analysis of dag-consistent distributed shared-memory algorithms. In *Proceedings of the 8th Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA'96)*. ACM, New York, NY, 297–308. DOI: <http://dx.doi.org/10.1145/237502.237574>
- Robert D. Blumofe, Christopher F. Joerg, Bradley C. Kuszmaul, Charles E. Leiserson, Keith H. Randall, and Yuli Zhou. 1995. Cilk: An efficient multithreaded runtime system. In *Proceedings of the 5th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP'95)*. ACM, New York, NY, 207–216. DOI: <http://dx.doi.org/10.1145/209936.209958>
- Robert D. Blumofe and Charles E. Leiserson. 1998. Space-efficient scheduling of multithreaded computations. *SIAM Journal on Computing* 27, 1, 202–229. DOI: <http://dx.doi.org/10.1137/S0097539793259471>
- Robert D. Blumofe and Charles E. Leiserson. 1999. Scheduling multithreaded computations by work stealing. *Journal of the ACM* 46, 5, 720–748. DOI: <http://dx.doi.org/10.1145/324133.324234>
- F. Warren Burton and M. Ronan Sleep. 1981. Executing functional programs on a virtual tree of processors. In *Proceedings of the 1981 Conference on Functional Programming Languages and Computer Architecture (FPCA'81)*. ACM, New York, NY, 187–194. DOI: <http://dx.doi.org/10.1145/800223.806778>
- David Chase and Yossi Lev. 2005. Dynamic circular work-stealing deque. In *Proceedings of the 17th Annual ACM Symposium on Parallelism in Algorithms and Architectures (SPAA'05)*. ACM, New York, NY, 21–28. DOI: <http://dx.doi.org/10.1145/1073970.1073974>
- Matthew Fluet, Mike Rainey, John Reppy, and Adam Shaw. 2008. Implicitly-threaded parallelism in manticore. In *Proceedings of the 13th ACM SIGPLAN International Conference on Functional Programming (ICFP'08)*. ACM, New York, NY, 119–130. DOI: <http://dx.doi.org/10.1145/1411204.1411224>
- Matteo Frigo, Charles E. Leiserson, and Keith H. Randall. 1998. The implementation of the Cilk-5 multithreaded language. In *Proceedings of the ACM SIGPLAN 1998 Conference on Programming Language Design and Implementation (PLDI'98)*. ACM, New York, NY, 212–223. DOI: <http://dx.doi.org/10.1145/277650.277725>
- John Giacomoni, Tipp Moseley, and Manish Vachharajani. 2008. FastForward for efficient pipeline parallelism: A cache-optimized concurrent lock-free queue. In *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP'08)*. ACM, New York, NY, 43–52. DOI: <http://dx.doi.org/10.1145/1345206.1345215>
- Michael I. Gordon, William Thies, and Saman Amarasinghe. 2006. Exploiting coarse-grained task, data, and pipeline parallelism in stream programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS XII)*. ACM, New York, NY, 151–162. DOI: <http://dx.doi.org/10.1145/1168857.1168877>
- Robert H. Halstead, Jr. 1984. Implementation of multilisp: Lisp on a multiprocessor. In *Proceedings of the 1984 ACM Symposium on LISP and Functional Programming (LFP'84)*. ACM, New York, NY, 9–17. DOI: <http://dx.doi.org/10.1145/800055.802017>
- Robert H. Halstead, Jr. 1985. MULTILISP: A language for concurrent symbolic computation. *ACM Transactions on Programming Languages and Systems* 7, 4, 501–538. DOI: <http://dx.doi.org/10.1145/4472.4478>

- D. A. Kranz, R. H. Halstead, Jr., and E. Mohr. 1989. Mul-T: A high-performance parallel Lisp. In *Proceedings of the ACM SIGPLAN 1989 Conference on Programming Language Design and Implementation (PLDI'89)*. ACM, New York, NY, 81–90. DOI: <http://dx.doi.org/10.1145/73141.74825>
- I-Ting Angelina Lee, Charles E. Leiserson, Tao B. Schardl, Jim Sukha, and Zhunping Zhang. 2013. On-the-fly pipeline parallelism. In *Proceedings of the 25th ACM Symposium on Parallelism in Algorithms and Architectures (SPAA'13)*. ACM, New York, NY, 140–151. DOI: <http://dx.doi.org/10.1145/2486159.2486174>
- Daniel Spoonhower, Guy E. Blelloch, Phillip B. Gibbons, and Robert Harper. 2009. Beyond nested parallelism: Tight bounds on work-stealing overheads for parallel futures. In *Proceedings of the 21st Annual Symposium on Parallelism in Algorithms and Architectures (SPAA'09)*. ACM, New York, NY, 91–100. DOI: <http://dx.doi.org/10.1145/1583991.1584019>

Received December 2014; revised May 2015; accepted June 2015